

# Part 1: Introduction — What Is a Docker Image?

A **Docker image** is (more than) a **blueprint** for a container. It includes the actual state executed more than the blueprint. It's a **read-only, layered filesystem** built using a **Dockerfile**.

When you run an image with `docker run`, Docker creates a **container** — a running instance with a writable layer on top of the image.

---

# Part 2: Dockerfile Fundamentals

A **Dockerfile** is a **recipe** that defines how your image is built.

Here's a minimal example:

```
# Simple Dockerfile Example
FROM python:3.12-slim
WORKDIR /app
COPY . .
RUN pip install -r requirements.txt
CMD ["python", "app.py"]
```

Let's break down **each instruction**.

---

## 1. FROM — Base Image

**Syntax:**

```
FROM <image>[ :<tag>]
```

- It specifies the **base image** (starting point).
- Every image must start with a `FROM` (except for scratch images).

`ARG` is the only command that can be before `FROM`. But that is rarely used.

**Examples:**

```
FROM ubuntu:22.04
FROM python:3.12-slim
FROM node:20-alpine
```

**Best Practices**

- Use **lightweight images** (`-alpine`, `-slim`) to reduce image size.
  - Always **pin versions** (e.g., `python:3.12-slim`) for reproducibility.
  - Use **scratch** for building minimal images (for Go, Rust, etc.).
- 

## 2. RUN — Execute Commands During Build

**Syntax:**

```
RUN <command>
RUN ["executable", "param1", "param2"]
```

Used to install software, set up environment, or configure files — it creates a **new image layer**.

**Example:**

```
RUN apt-get update && apt-get install -y curl
RUN pip install flask
```

### Best Practices

- Combine related commands to reduce image layers:

```
RUN apt-get update && apt-get install -y curl python3-pip && rm -rf /var/lib/apt/lists/*
```

- Always **clean up caches** and temporary files.

- Use **multi-stage builds** for smaller final images.

---

## 3. CMD — Default Command to Run in Container

### Syntax:

```
CMD ["executable", "param1", "param2"] # exec form (recommended)
CMD command param1 param2           # shell form
```

- Defines the **default command** that runs when the container starts.
- You can **override it** using `docker run <image> <your_command>`.

### Example:

```
CMD ["python", "app.py"]
```

### Best Practices

- Always use **JSON (exec) form** — avoids spawning an extra shell.
- Only **one CMD** per Dockerfile — the last one overrides others.

---

## 4. ENTRYPOINT — Fixed Main Command

### Syntax:

```
ENTRYPOINT ["executable", "param1", "param2"]
```

- Used to define the **main application** that should always run.
- CMD is often used to provide **default arguments** to ENTRYPOINT.

### Example:

```
ENTRYPOINT ["python", "app.py"]
```

Now even if you run:

```
docker run myapp arg1
```

it executes:

```
python app.py arg1
```

### Best Practice

Use ENTRYPOINT for your main app and CMD for its arguments.

---

## 5. Combining ENTRYPOINT and CMD

You can use both together for flexibility.

### Example:

```
FROM ubuntu:22.04
ENTRYPOINT ["echo"]
CMD ["Hello, World!"]
```

If you run:

```
docker run myimage
```

Output → Hello, World!

If you override CMD:

```
docker run myimage Goodbye
```

Output → Goodbye

#### Concept:

- `ENTRYPOINT` = Fixed executable.
  - `CMD` = Default argument.
- 

## 6. ADD and COPY — Copying Files into the Image

Both copy files from **host → image**, but they differ slightly.

### COPY

Simplest and most common — just copies files/directories.

```
COPY requirements.txt /app/  
COPY . /app
```

**Best Practice:** Use `COPY` whenever possible — it's predictable.

---

### ADD

Adds files **and** supports:

- **Remote URLs**
- **Automatic tar extraction**

```
ADD https://example.com/file.tar.gz /tmp/  
ADD myapp.tar.gz /app/
```

**Best Practice:** Use only when you need those extra features.  
Don't use `ADD` for normal local file copies — it's less transparent.

---

## 7. ENV — Set Environment Variables

**Syntax:**

```
ENV <key>=<value>
```

**Example:**

```
ENV APP_ENV=production  
ENV PORT=8080
```

Used to configure runtime environment for your app.

#### Best Practices

- Use `ENV` for config that shouldn't change often.
  - Use `docker run -e KEY=value` for dynamic values.
- 

## Part 3: Building a Real-World Docker Image Example

Let's build a **Flask web app** image — production ready.

---

### Directory structure:

```
flask-app/
├── app.py
├── requirements.txt
└── Dockerfile
```

#### app.py

```
from flask import Flask
app = Flask(__name__)

@app.route('/')
def hello():
    return "Hello from Docker!"

if __name__ == '__main__':
    app.run(host='0.0.0.0', port=8080)
```

#### requirements.txt

```
flask==3.0.2
```

---

#### Dockerfile

```
# Stage 1: Base
FROM python:3.12-slim

# Set working directory
WORKDIR /app

# Copy dependency file first for caching
COPY requirements.txt .

# Install dependencies
RUN pip install --no-cache-dir -r requirements.txt

# Copy the rest of the source code
COPY . .

# Set environment variables
ENV PORT=8080

# Expose port
EXPOSE 8080

# Use entrypoint and cmd
ENTRYPOINT [ "python" ]
CMD [ "app.py" ]
```

---

### Build and Run

```
docker build -t flask-app:latest .
docker run -p 8080:8080 flask-app
```

Access at → <http://localhost:8080> ↗

---

## Part 4: Docker Image Best Practices

Category	Best Practice	Why
Base Image	Use minimal images (-alpine, -slim)	Reduces size and attack surface
Layers	Combine related RUNs	Fewer layers = smaller image

<b>Caching</b>	Copy dependency files early	Avoids reinstalling packages unnecessarily
<b>Cleanup</b>	Remove caches ( <code>rm -rf /var/lib/apt/lists/*</code> )	Prevents bloat
<b>Security</b>	Use non-root user ( <code>USER appuser</code> )	Prevents privilege escalation
<b>Scanning</b>	Use docker scan or Trivy	Detects vulnerabilities
<b>Secrets</b>	Don't hardcode passwords or tokens	Use environment variables or Docker secrets
<b>Versioning</b>	Pin base images and dependencies	Ensures reproducibility
<b>Multi-stage Builds</b>	Use multiple stages to reduce size	Only final stage goes to production
<b>Healthcheck</b>	Add <code>HEALTHCHECK</code> instruction	Detect and auto-restart unhealthy containers

## Example — Secure Image Setup

```
FROM python:3.12-slim
WORKDIR /app

COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt && \
    useradd -m appuser

COPY . .
USER appuser
EXPOSE 8080
ENTRYPOINT [ "python" ]
CMD [ "app.py" ]
```

Runs as non-root  
Minimal image  
Cached dependency installation

## Part 5: Bonus — Inspecting and Analyzing Your Image

- List image layers:

```
docker history flask-app
```

- Inspect metadata:

```
docker inspect flask-app
```

- Scan for vulnerabilities:

```
docker scan flask-app
```

or use **Trivy**:

```
trivy image flask-app
```

Detailed reference:

1. Refer to my repository in
  - a. <https://github.com/vilasvarghese/docker-k8s/dockerfiles>
2. Command to use
  - a. `docker build -t <image name in lower case> -f <file name>` .