

Docker containers

Agenda

What exactly is a container - a deep dive

1. Why Containers Are Ephemeral

A container is simply:

a Linux process + a temporary root filesystem.

The filesystem layers (image layers + writable layer) exist only while the container runs.

Key Point:

If the main container process exits → container stops → writable layer disappears → the container is gone.

This makes containers *stateless* and *ephemeral by design*.

2. Container = Just a Process on the Host


When you run:

```
docker run nginx
```

Docker does NOT start a virtual machine.

It launches a **Linux process on the host**, but inside kernel namespaces.

Diagram



The "container" is simply the **nginx process** running under isolation.

3. Why Container Isolation = Pseudo-Isolation

Namespaces hide parts of the world

Cgroups limit resource usage

BUT:

- All containers share the **same host kernel**
- All container processes appear in the host's `ps -ef`
- A root user inside the container **is not a real root** (User Namespace)
- A process inside a container can't break out *normally*, but under kernel vulnerabilities it **can escape**, because it's not a VM.

Thus:

Containers ≠ Virtual Machines

Containers = Regular Processes with illusions

That's why we call it **process-level isolation**, not OS-level isolation.

4. Understanding containerd, shim, runc

Modern container architecture:

`docker --> containerd --> shim --> runc --> your process`

containerd

- Long-running daemon
- Manages containers
- Handles images, snapshots, networking

containerd-shim


- One shim **per running container**
- Acts as parent for container process
- Allows containerd to restart without killing containers

runc

- The actual runtime

- Uses Linux kernel syscalls (`clone`, `unshare`, `pivot_root`)
- Sets up namespaces & cgroups
- Starts the main process inside the new isolated environment

Diagram



5. Process Tree Explains Container Behavior

Containers behave like this:

- A container is started by **containerd-shim**
- Shim launches **runc**
- runc sets up namespaces
- runc starts the container's *main process*, e.g., `/usr/sbin/nginx`

Process Tree Example

```

$ pstree -pl
containerd(233)
└─ containerd-shim(9421)
  
```

```
└─ nginx(9422)
    └─ nginx worker(9423)
        └─ nginx worker(9424)
```

IMPORTANT:

- The **main process inside the container** becomes **PID 1** *in the container's PID namespace*.
 - But on the host, it is just another process (like 9422).
-

6. Why killing the main process stops the container

Because:

The container lifecycle = lifecycle of its PID 1

If PID 1 exits in the container namespace:

- containerd-shim detects process exit
- shim notifies containerd
- containerd marks container as stopped
- writable layer is discarded
- networking is removed
- cgroups cleaned
- container disappears

Example:

```
docker kill <container>
```

This sends SIGKILL to PID 1.

Container stops instantly.

7. Why containers are ephemeral (with diagrams)

Each container has:

- a **read-only image**
- a **temporary writable layer**

Diagram

```
Image Layers (read-only)
  layer 1
  layer 2
  layer 3
-----
Container Writable Layer (deleted on stop)
  /var/log/app.log
  /tmp/files
-----
Running Process
```

When container stops → writable layer is deleted → all state lost.

Thus:

Container data disappears unless stored in:

- Volumes
 - Bind mounts
 - External storage (DB, S3, etc.)
-

8. Why containers appear like isolated machines

Namespaces create illusions:

PID Namespace

Container process sees itself as PID 1.

Network Namespace

Container sees its own:

- eth0
- IP address
- routing table

Mount Namespace

Container sees its own root filesystem.

UTS Namespace

Container sees its own hostname.

IPC Namespace

Own shared memory.

User Namespace

Root inside container \neq root on host.

9. Why containers start so fast

Because:

- No kernel boot
- No virtual hardware
- No BIOS
- No init system unless you add one

Containers start a single process:

runc \rightarrow clone syscall \rightarrow process runs

Start time: **$\sim 50\text{ms}$**

VMs take: **tens of seconds**

10. Why containers are light

Containers share:

- The host kernel
- The host OS
- Libraries (optional)
- CPU scheduler
- Memory management

Only the filesystem + process isolation is unique.

11. Why containers fail if the main process exits

Because a container = **one main process**.

Examples:

- nginx dies → container dies
- python app.py exits → container exits
- sleep 5 completes → container exits immediately

This confuses beginners.

Containers need:

- process managers (supervisord)
 - or multi-process apps designed properly
-

12. Summary in 10 bullet points

1. A container is just a **process + isolation**
 2. It runs on the **host OS**, not in a VM
 3. Isolation is via **namespaces**
 4. Resource limits via **cgroups**
 5. Process started by **runc**, managed by **containerd-shim**
 6. The container's root process is **PID 1** in its namespace
 7. If PID 1 dies → container stops
 8. Writable layer is temporary → ephemeral
 9. Isolation is not strong like a VM → "pseudo-isolation"
 10. All container processes visible in host `ps -ef`
-

CONTAINERS EXPLAINED — FROM ZERO TO DEEP INTERNALS

1. First: What EXACTLY is a container?

A **container** is NOT:

- a virtual machine
- a mini-OS
- something magical

A **container** IS:

A normal Linux process with special isolation applied using kernel features.

That's it.

Why does it *feel* like a separate computer?

Because:

- It has its own user space
- It has its own Inter process communicaton space
- It has its **own filesystem**
- Its **own network (space) interface**
- Its **own hostname**
- Its **own process tree**
- Its **own resource limits**

But none of this is done by creating a new OS.

All of it is done using:

- ✓ **Linux Namespaces**
 - ✓ **Linux Cgroups**
 - ✓ **Chroot / Filesystem overlays**
 - ✓ **Capabilities**
 - ✓ **Seccomp (optional)**
-

2. Why a container is just a process

Normally, when you run a process:

```
$ python app.py
```

the process shares:

- host network interfaces (eth0)
- host filesystem
- host PID number space
- host users
- host resources

In a container:

```
docker run python:3.11
```

Docker creates a new process:

- an instance of containerd-shim-runc-v2

python (inside container) → but still a Linux process on host

But before launching it, Docker enables isolation:

- ✓ PID namespace
- ✓ NET namespace
- ✓ UTS namespace
- ✓ MOUNT namespace
- ✓ USER namespace
- ✓ IPC namespace

Because of these namespaces, the process sees a **fake world**.

3. The Magic Behind Containers: Linux Namespaces

Think of namespaces like **AR (Augmented Reality) goggles**.

Each namespace puts a filter on what the process can see.

Let's break down the namespaces.

3.1 PID Namespace (Process ID Isolation)

Without namespace:

```
Host Processes:  
1 systemd  
22 sshd  
300 nginx  
4212 python
```


Inside a container:

```
/ # ps  
PID  Command  
1  python  
7  bash  
12 app
```

The container sees its own process as **PID 1**.

This makes the container feel like its own operating system.

Drawing:




3.2 Network Namespace

Each container gets its own:

- network stack
- IP address
- routing table
- firewall rules

But the host only sees a **veth pair**.

Drawing:



This is why containers have **their own IPs**.

3.3 Mount Namespace (Filesystem Isolation)

Each container sees **its own root filesystem (/)**.


But under the hood, Docker is doing:

- OverlayFS layers
- Chroot jail
- Bind mounts for volumes

Host filesystem:



Container filesystem:



Different root = **feels like a different machine**.

3.4 UTS Namespace (Hostname Isolation)

Lets a container have its own:

- hostname

- domain name

So inside container:

```
hostname = web-app
```

But host's hostname is:

```
hostname = ip-10-0-0-12
```

3.5 IPC Namespace (Shared Memory Isolation)

Containers get their own:

- semaphores
 - shared memory segments
-

3.6 USER Namespace

Allows containers to run as:

- root **inside the container**
 - but mapped to a non-root UID on the host
-

4. CGroups – Controlling Resources

Namespaces = isolation

Cgroups = resource limits

Cgroups limit:

- CPU
- Memory
- Disk I/O
- PIDs


Example:

```
Limit container to 256MB RAM
Limit container to 0.5 CPU
```

If you run too many processes:

cgroups prevents container from using more than allowed

Drawing:



5. So How Does Docker Create a Container?

Step-by-step:

1. Pull an image
2. Unpack filesystem layers (OverlayFS)
3. Create cgroups for limits
4. Create namespaces
5. chroot into the new root filesystem
6. Start the process (e.g., python)

So docker run:

```
docker run nginx
```

is actually:

```
create-net-namespace
create-mount-namespace
create-pid-namespace
apply-cgroups
chroot to overlayfs
launch /usr/sbin/nginx
```

6. Why your normal programs DO NOT feel like containers


Your normal programs:

- share host network
- share host filesystem
- share host hostname
- share host PID tree
- share all CPU/RAM


Containers **hide all of that.**

7. VM vs Container (Deep Difference)

VM Architecture



Container Architecture



Docker does NOT run a separate kernel → only processes.