

# ReplicaSet – Deep Technical Tutorial

A **ReplicaSet (RS)** is a Kubernetes controller responsible for ensuring a **specified number of Pod replicas** are always running.

A ReplicaSet **continuously compares**:

```
desired state (spec.replicas)  
vs  
current state (number of pods matching label selector)
```

If mismatch:

- It **creates** Pods.
- It **deletes** excess Pods.
- It **replaces** failed Pods.

It does this via the **Kubernetes control plane**, mostly through:

- kube-apiserver
  - kube-controller-manager
  - kube-scheduler
  - kubelet
  - etcd
- 
- 

## 1. Let's Start With a YAML

```
apiVersion: apps/v1  
kind: ReplicaSet  
metadata:  
  name: nginx-rs  
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      app: web  
  template:  
    metadata:  
      labels:  
        app: web
```

```
spec:  
  containers:  
    - name: nginx  
      image: nginx:latest
```

You run:

```
kubectl apply -f nginx-rs.yaml
```

---

## ----- 2. What Happens? — Full Control Plane Flow

### Step 1: kubectl → API Server

kubectl sends the ReplicaSet definition as a REST request to:

```
kube-apiserver:443/apis/apps/v1/replicasets
```

The API server:

1. Authenticates (AuthN)
2. Authorizes (RBAC)
3. Admits (Validating/Mutating Webhooks)
4. Stores the ReplicaSet object in etcd

#### etcd state after storing ReplicaSet:

```
/registry/replicasets/default/nginx-rs  
  - spec.replicas = 3  
  - selector = app=web  
  - pod template stored here
```

---

## ----- 3. ReplicaSet Controller Starts Acting

Running inside **kube-controller-manager**, the **replicaset controller** watches the API server via a watch stream.

It receives an event:

```
ADDED ReplicaSet "nginx-rs"
```

### The controller compares:

```
current pod count = 0  
desired pod count = 3
```

So it decides to **create 3 Pods**.

### It sends Pod objects to the API server:

```
POST /api/v1/namespaces/default/pods
```

Each Pod uses the template stored inside the ReplicaSet.

---

## 4. etcd After Pod Creation

Three pod entries appear:

```
/registry/pods/default/web-abc12  
/registry/pods/default/web-def34  
/registry/pods/default/web-ghi56
```

Within each Pod object:

- `metadata.ownerReferences` points to the ReplicaSet
  - `status.phase = Pending`
  - `spec.nodeName = ""` (no node assigned yet)
- 

## 5. Scheduler Now Gets Involved

The **kube-scheduler** continuously watches the API server for **Pending** pods with no node assigned.

It receives:

```
ADDED Pod web-abc12  
ADDED Pod web-def34
```

```
ADDED Pod web-ghi56
```

For each pod:

1. Scheduler runs its scheduling logic (fit + priority).
2. Selects a node (example: node1).
3. Writes **binding** back to API server:

```
POST /api/v1/namespaces/default/pods/web-abc12/binding
spec.nodeName = node1
```

### **etcd after scheduling:**

Each pod gets updated:

```
spec.nodeName = node1 (or node2, node3 ... )
status.phase = Scheduled (not Running yet)
```

---

-----

## **6. Kubelet on Each Node Takes Over**

Each worker node's **kubelet** watches:

```
/api/v1/pods?fieldSelector=spec.nodeName=<this node>
```

When kubelet sees the Pod assigned to its node:

1. It pulls the image
2. Creates container via container runtime (Docker, containerd, CRI-O)
3. Starts it
4. Updates the API server with pod status:

```
status.phase = Running
```

**This updated status is stored in etcd.**

---

-----

## **ReplicaSet Controller Continuous**

# Reconciliation Loop

The ReplicaSet controller ALWAYS runs this loop:

## Read from API server:

```
List Pods matching selector app=web
```

Compare:

- desired: 3
- current: 3

Equal → do nothing.

---

## Failure Scenarios — How ReplicaSet Reacts

### Pod Crashes

kubelet updates:

```
status.phase = Failed
```

ReplicaSet sees:

```
desired = 3  
running = 2
```

→ Creates a new Pod object.

### Node Dies

If node stays **NotReady** for ~5 minutes, node controller marks all pods as:

```
status.phase = Unknown
```

Then the ReplicaSet controller deletes those pods and creates replacements.

### You manually delete a Pod

If deletion is not with --cascade=orphan, the RS sees:

```
2 running  
desired 3
```

→ Creates 1 new pod.

---



## If You Delete the ReplicaSet Itself

You run:

```
kubectl delete rs nginx-rs
```

The API server marks RS for deletion → stored in etcd.

The ReplicaSet controller receives:

```
DELETED ReplicaSet nginx-rs
```

It deletes all **owned pods**.

All pod entries are removed from etcd.

---

## Putting It All Together — ETCD ILLUSTRATION

Below is a simplified view of how objects exist in etcd:

```
/registry  
  -- replicaset  
    -- default  
      -- nginx-rs  
        - spec.replicas = 3  
        - selector: app=web  
        - podTemplate  
  
  -- pods  
    -- default  
      -- web-abc12  
        - ownerRef = nginx-rs  
        - spec.nodeName = node1  
        - status: Running  
      -- web-def34
```

```
    └── web-ghi56  
        └── nodes  
            └── node1  
            └── node2  
            └── node3
```

---

```
-----
```

## ReplicaSet Summary (Deep Insight)

Component	Role
<b>kubectl</b>	Sends ReplicaSet YAML to API server
<b>API Server</b>	Validates and stores object in etcd
<b>etcd</b>	Stores desired state for RS and Pods; stores actual pod status
<b>ReplicaSet Controller</b>	Detects pod count mismatch and creates/deletes pods
<b>Scheduler</b>	Assigns pods to nodes
<b>Kubelet</b>	Creates, manages containers on node; reports status
<b>Container Runtime</b>	CRI-O/Docker/containerd actually runs containers

ReplicaSets are purely **controllers** — they never run pods; they only ensure correct count.